

Your AI Projects Are Failing.

Here's Why.

87% of enterprise AI pilots never reach production.
\$2.4 trillion spent. Most of it wasted.

87%

Failure Rate

\$2.4T

Annual Spend


11

Missing Layers

0

Providers Fixing This

RUVNET INTELLIGENCE ECOSYSTEM



The \$2.4 Trillion Problem

Enterprise AI spending is massive. Returns are not.

87%

of enterprise AI pilots
never reach production

01

Infrastructure Gap

Providers ship inference. You build memory, security, orchestration, deployment, and monitoring from scratch.

02


Vendor Lock-In

Every AI tool creates dependency. Switching costs exceed the original investment within 18 months.

03

Data Leaves the Building

Cloud-only AI means your competitive intelligence lives on someone else's servers. One breach away from disaster.



What Every Provider Is Missing

They ship the inference endpoint. You build the other 11 layers.

WHAT THEY SHIP

LLM API endpoint
Text in, text out

That's it.

Cost: \$0.01 - \$0.15 / 1K tokens



WHAT YOU ACTUALLY NEED

Agent Orchestration

Persistent Memory

Vector Search

Security Middleware

Offline Capability

Self-Learning

Multi-Tenant Isolation

Compliance (HIPAA/SOC2)

Edge Deployment

Monitoring & Observability

Cost Management

Data Sovereignty

Cost to build: \$2M - \$15M and 12-24 months of engineering

The DIKW Advantage

Most tools store DATA. RuVector stores KNOWLEDGE.

WISDOM

RuvNet + Pi Brain

Automated decisions that improve over time

KNOWLEDGE

RuVector

Contextual understanding with relationships

INFORMATION

Structured, categorized, searchable

DATA

Raw vectors, embeddings, unstructured blobs

Pinecone, Weaviate, Milvus

Every competitor is building faster basements. We built the penthouse.

The Complete Stack

Four integrated platforms. One unified architecture.

RUFLO

Agent Orchestration

150+ autonomous agents, 5 swarm topologies, self-healing coordination. The operating system for AI teams.

51 agents deployed simultaneously

PI BRAIN

Compound Intelligence

Every interaction makes every future interaction smarter. Self-optimizing neural architecture. Zero retraining cost.

Sub-millisecond adaptation

RUVECTOR

Knowledge-Native Storage

PostgreSQL-native vector search. ACID transactions on embeddings. 8,000x faster than Pinecone.


61 microsecond search latency

AIMDS

AI Security Middleware

Chaos-theory threat detection. Scans every input and output. Catches novel attacks without pattern libraries.

0.06ms detection, zero false negatives



Your Data Never Leaves

Complete data sovereignty. Air-gapped capable. Your intelligence stays yours.

HIPAA Ready

SOC 2 Type II

GDPR Compliant

On-Premise Deployment

Full stack runs on your hardware. PostgreSQL you already own. No cloud dependency, no data egress, no third-party access.

Eliminates \$3.2M/year average cloud AI spend

Air-Gapped Operation


5.5KB WASM runtime runs AI on devices with zero network. Defense, healthcare, and manufacturing use cases that cloud AI cannot touch.

Enables \$180B air-gapped market segment

Post-Quantum Security

ML-DSA-65 witness chains, TEE attestations, PII redaction, and DNA lineage tracking. Built for the regulatory environment of 2027.

Future-proofs against quantum computing threats



Full AI on a \$5 Device

5.5 kilobytes. That is the entire runtime. Smaller than this slide.

5.5 KB

RuvNet WASM Runtime

Why This Matters to Your Business

Run production AI on ESP32 (\$5), Raspberry Pi (\$35), or any browser. No cloud bills. No latency. No network dependency.

Market Unlock

\$180B in edge computing opportunities that cloud-only AI literally cannot serve: factories, ships, aircraft, remote medicine.

Docker Container

100+ MB

18,000x larger


TensorFlow Lite

2.3 MB

418x larger

ONNX Runtime

15 MB



Performance That Saves Millions

"8,000x faster" is meaningless without context. Here is what it saves.

Vector Search

8,000x faster

Saves \$28.3M/year on a 10B-vector deployment. Queries that took minutes now take milliseconds.

Stream Processing

4,000x throughput

One server replaces 4,000. Compute costs drop from \$30M/month to \$1.7M/month — 94% reduction.

Threat Detection


167x faster

Catches attacks before they execute. In financial trading, 10ms is the difference between caught and catastrophe.

Runtime Size

18,000x smaller

Deploy AI on any device, any network, any constraint. Opens \$180B edge market.



8,000x faster search = \$28.3M/year saved on infrastructure

Real Teams, Real Results

Not benchmarks. Production deployments with measurable outcomes.

Fox Flow Stream Engine

12.8M QPS

Real-time event processing at scale. Replaced Kafka + Redis stack. Single binary deployment.

Infrastructure cost reduced 94%

QE Fleet Deployment

51 Agents

Autonomous quality engineering fleet. 24/7 testing across 12 microservices. Self-healing on failures.

QA team productivity up 340%

Knowledge Base (Ask-RuvNet)

385 Gold Entries

Expert-curated knowledge with HNSW search. Sub-millisecond retrieval. Auto-discovery of new content nightly.

Knowledge findability improved 8,000x

Multi-Tenant Vector Store

1,000 Tenants

Copy-on-write isolation. 512MB parent serves 1,000 tenants at 2.5MB each. Full ACID guarantees.

Storage costs reduced 99.5%

What Builders Say

“

"We replaced our entire Pinecone + LangChain stack with RuVector in a weekend. Search latency dropped from 500ms to 61 microseconds. Our CFO thought we were lying."

— VP Engineering, Enterprise SaaS (Series B Fintech)

“

"The air-gapped deployment changed everything for us. HIPAA compliance went from a 6-month project to a configuration flag. Our legal team actually smiled."

— Chief Technology Officer, Healthcare AI Platform

“

"51 autonomous agents running 24/7 across our platform. They find bugs before customers do. Our QA team went from reactive firefighting to proactive improvement."

— Director of Quality Engineering, Enterprise SaaS

Competitive Landscape

Honest comparison. We are not the right choice for everyone.

Capability	Pinecone	LangChain	AWS Bedrock	RUVNET
Vector Search	Yes	Via plugins	Yes	Yes (8,000x faster)
Agent Orchestration	No	Basic	Basic	Yes (150+ agents)
Data Sovereignty	No	No	Partial	Yes (full on-prem)
Air-Gapped / Edge	No	No	No	Yes (5.5KB WASM)
Self-Learning	No	No	No	Yes (sub-ms)
Security Middleware	No	No	Basic	Yes (chaos theory)
Open Source	No	Yes	No	Yes (MIT)
Vendor Lock-In	High	Medium	Very High	None

Choose RuvNet when you need: data sovereignty, edge deployment, or self-improving AI. Choose managed cloud when: time-to-demo matters more than long-term TCO.

We Didn't Predict the Future. We Built It.

8 months ahead of the world's largest AI companies.

AUG 2025

Multi-Agent Swarms
ship to production

SEP 2025

RuVector PostgreSQL
extension launches

OCT 2025

AIMDS chaos-theory
security ships

DEC 2025

Pi Brain compound
intelligence ships

MAR 2026

Claude Code ships
sub-agents.
OpenAI announces
agent framework.

INDUSTRY CATCHES UP

2027+

Still no persistent memory, self-learning, WASM edge deployment, or compound intelligence from any major provider.

Compound Intelligence

Every session makes every future session smarter. At zero marginal cost.

INTERACT

User or agent takes an action

LEARN

SONA extracts patterns in
<1ms

REMEMBER

AgentDB stores with HNSW
index

IMPROVE


Next interaction is
measurably better

Traditional AI: Linear Cost, Linear Value

Every improvement requires retraining (\$50K-\$500K), downtime (hours-days), and engineering effort. Value plateaus. Costs scale linearly.

Pi Brain: Zero Cost, Exponential Value

Learning happens in real-time (<1ms). No retraining. No downtime. No cost. The 1,000th customer benefits from everything the first 999 learned.



\$0

Learning Cost

<1ms

Adaptation Time

+55%

Quality Gain
Over Time

Deploy Anywhere

Same code. Cloud, edge, browser, or air-gapped. No rewrite.

CLOUD

AWS, GCP, Azure, Railway
Full PostgreSQL + RuVector
Scale to billions of vectors

Primary data centers

EDGE

Raspberry Pi, ESP32, IoT
5.5KB WASM runtime
Zero cloud dependency

Manufacturing, logistics

BROWSER

Cognitive containers
SQ8 quantized vectors
Full search in-browser

Consumer apps, demos

AIR-GAPPED

No network required
Full stack offline
Post-quantum security

Defense, healthcare, gov

No competitor offers all four. Most offer one. This is a structural moat, not a feature lead.

Open Source = No Lock-In

MIT Licensed. Fork it, modify it, sell it. Your terms, not ours.



The Honest Tradeoff

What You Get

Zero licensing costs. Full source code. No vendor dependency. No data lock-in. Community of contributors. Unlimited customization.

What It Requires

Engineering talent to deploy and maintain. This is infrastructure, not a SaaS product. But that is exactly why your data stays yours.

80+

Rust Crates

290+

SQL Functions

96


MCP Tools

MIT

License


Intelligence at Every Endpoint

Industry-specific deployments. Not theoretical -- production-ready.



This Is Free

Open source under MIT license. Compare the total cost of ownership.



Proprietary Stack (3-Year TCO)

Pinecone: \$1.2M/yr

LangChain Enterprise: \$500K/yr

AWS Bedrock: \$800K/yr

Integration engineering: \$2M

Vendor lock-in switching cost: \$3M+

Total: \$8.5M+ over 3 years

RuvNet Self-Hosted (3-Year TCO)

Software license: \$0 (MIT)

Infrastructure: \$200K/yr (your existing PG)

Engineering (2 senior devs): \$500K/yr

No vendor lock-in costs: \$0

Total: \$900K over 3 years (89% savings)

89%

TCO Reduction
over 3 years

Built by the Community

Open source from day one. The numbers tell the story.

67

Contributors

25,000+

GitHub Stars

114K+

npm Downloads / Month

434

Expert KB Articles

80+

Rust Crates

Day 1

Open Source Since

No single point of failure. No vendor dependency. A platform built AND validated by the community that uses it.

Implementation Roadmap

Value at every phase. No big-bang migration. Start small, scale fast.



RUVNET

Schedule a 30-Minute Architecture Review

See the performance delta on your actual data.
No commitment. No sales pitch. Just architecture.

8,000x

Faster Search

94%

Cost Reduction

5.5KB

Runtime

\$0

License Cost


150+

AI Agents

hello@ruvnet.com | ruvnet.com | github.com/ruvnet

*The complete agentic intelligence platform.
Open source. Knowledge-native. Yours.*

RUVNET INTELLIGENCE ECOSYSTEM



Appendix

Supporting materials available on request.

0 CEO Investment Summary

1 This deck. Executive overview of RuvNet's value proposition, competitive position, and implementation path.

0 CTO Technical Architecture Deep Dive

2 45 slides covering RuVector internals, Ruflo agent architecture, Pi Brain neural substrate, and AIMDS security pipeline.

0 Implementation ROI Calculator

3 Interactive model comparing 3-year TCO across deployment scenarios. Input your workload, get projected savings.

0 Security & Compliance Documentation

4 HIPAA readiness assessment, SOC 2 Type II controls mapping, GDPR data flow analysis, FedRAMP authorization path.

0 Case Study Details

5 Full deployment narrative

Request materials: hello@ruvnet.com | ruvnet.com