

R U V N E T

Technical Architecture Deep Dive

148 capabilities across 14 modules. Production-grade agentic infrastructure.



80+

Rust Crates

290+

SQL Functions

39

Attention Mechanisms

265K

Lines of Code


What Providers Actually Ship

They ship the inference endpoint. You build the other 11 layers.

Capability	OpenAI	Anthropic	Google	AWS	RuvNet
LLM API	✓	✓	✓	✓	✓
Multi-Agent (150+ types)	✗	✗	✗	✗	✓
Persistent Memory (4 types)	✗	✗	✗	✗	✓
Vector Search (<100μs)	✗	✗	✗	✗	✓
Security Middleware	✗	✗	✗	✗	✓
Offline / Air-Gapped	✗	✗	✗	✗	✓
Self-Learning (<1ms)	✗	✗	✗	✗	✓
WASM Runtime (7.2KB)	✗	✗	✗	✗	✓
Anti-Hallucination (formal)	✗	✗	✗	✗	✓
Graph Analytics (GNN)	✗	✗	✗	✗	✓
MCP Protocol (96 tools)	✗	○	✗	✗	✓
Data Sovereignty	✗	✗	✗	○	✓

System Architecture

14 modules. Unified data plane. Security envelope.



RVF: 3-Tier Compute

WASM (5.5KB) | eBPF (kernel bypass) | Unikernel (125ms Linux boot)

Nervous System: 5 Bio-Inspired Layers

Sensing (<100ns) | Reflex (<1μs) | Memory (10^{30} cap) | Learning | Coherence

PostgreSQL Unified Data Plane — 290+ SQL Functions | HNSW Indexing | RuVector Extension | Drop-in pgvector replacement

Security Envelope — Witness Chains (ML-DSA-65 Post-Quantum) | TEE Attestations | PII Redaction | DNA Lineage Tracking

96 MCP Tools

80+ Rust Crates

49+ npm Packages

25+ ADRs

14 Modules

This entire stack is open source. MIT license. Free forever.

ONLY POSSIBLE WITH RUVNET

Enterprise AI Without Data Leakage

Every AI API sends your data to someone else's servers. RuVector runs inside your PostgreSQL.

Standard AI Architecture

Your Data → Internet → OpenAI/Google Servers
→ Response

- Data transmitted to third-party infrastructure
- No audit trail of what happens to your data
- Compliance teams block adoption
- Network dependency = single point of failure

RuvNet Architecture

Your Data → Your PostgreSQL → RuVector (in-process) → Response

- 290+ SQL functions run in your DB process
- AIMDS scans inbound AND outbound for leaks
- Witness chains: cryptographic audit of every op
- Air-gap with RVF containers — zero network code

EXAMPLE: Bank Loan Portfolio Analysis

10 years of loan performance data. Millions of records. With standard AI: compliance blocks the project because data would leave the network. With RuVector: deploy as a PostgreSQL extension on your existing DB. AIMDS enforces PII redaction. AI analyzes everything in-place. HNSW finds similar historical patterns in 61μs. Compliance signs off day one because zero data leaves the building.

ONLY POSSIBLE WITH RUVNET

Ship AI in 5.5 Kilobytes

RVF containers pack model + logic + vector search + memory + security into one file.
The WASM runtime is 5.5KB. That's smaller than this sentence in most fonts.

rvDNA: Genomics on a Phone

3B base pairs | <100ms | 7.2KB WASM

3 BILLION DNA base pairs processed in <100ms. Variant calling at 155 ns/SNP. Horvath Clock biological age analysis. Pharmacogenomics drug dosing. All in a browser tab. A \$50 phone becomes a genomics lab.

Offline Field Intelligence

Zero backend | Zero latency | Full AI

Entire equipment manual + troubleshooting history + diagnostic AI in one .rvf file. The file IS the server. Works underground, underwater, in a mine shaft. Agent carries full organizational knowledge without connectivity.

Browser-Native Enterprise Search

7.2KB engine | Client-side | Zero infra

Micro-HNSW runs the complete vector search engine client-side. User opens a web page — no installation, no server, no API keys. Search 100K+ documents semantically in the browser. IT deploys nothing.

Three compute tiers: WASM (5.5KB, browser) → eBPF (kernel bypass, 10Gbps+) → Unikernel (125ms Linux boot, bare metal)

ONLY POSSIBLE WITH RUVNET

Integrate at the Knowledge Level, Not SQL

SQL-Level Integration (Traditional)

- Months of ETL pipeline development
- Breaks when source schemas change
- Can't join unstructured data (PDFs, Slack, email)
- \$500K-\$2M projects, 6-12 months delivery

Knowledge-Level Integration (RuvNet)

- RuVector embeds everything into semantic space
- Self-healing: adapts automatically to changes
- Joins by meaning: PDFs + SQL + Slack + email
- Weeks, not months. Knowledge graphs, not ETL.

What This Enables

Microbiome Discovery

250M agents analyzed 40GB across 8,000 biological samples. GNN + HNSW discovered Huntington's disease ↔ gut microbiome correlations that human researchers hadn't identified. 187M interactions/sec. The AI was the researcher.

Enterprise-Wide Search

"Show me everything about the Acme deal" searches Salesforce, Jira, Confluence, Slack, email, and PDFs — by meaning, not keywords. One HNSW query, 61μs, every system. No connectors, no schema mapping.

Legacy COBOL Bridge

COBOL systems from the 1990s connected to modern REST APIs via knowledge embeddings in one afternoon. No rewrite. No translation layer. The vector embeddings understand both COBOL copybooks and JSON schemas.

Ruflo: Multi-Agent Orchestration

The brain that coordinates 150+ agent types across 5 network topologies.

Swarm Topologies

Hierarchical	Queen-led with anti-drift detection
Mesh	Peer-to-peer, Byzantine fault tolerant
Ring	Token-passing coordination
Star	Central hub routing
Adaptive	Topology shifts under load

Scale & Coordination

Hive Mind	1,000+ agents with sub-ms shared memory via PostgreSQL advisory locks
Emily OS	2 to 100,000 agents, LLM-agnostic, horizontal scaling
Raft Consensus	Leader election, multi-master replication with CRDT support
Gossip Protocol	Drift detection, behavioral monitoring across agents

150+

Agent Types

5

Topologies

100K

Agent Scale

<1ms

Shared Memory

RuVector: PostgreSQL-Native Vector Search

Not a separate vector database. A PostgreSQL extension. ACID transactions on vector operations.

Core Capabilities

- 290+ SQL functions in-process
- HNSW indexing (61μs search latency)
- Drop-in pgvector replacement
- ACID transactions on vector ops
- Joins vectors with relational data
- 80+ Rust crates for extensions
- Progressive indexing (L0/L1/L2)
- Temperature-tiered quantization

RVCOW: Multi-Tenant Economics

512MB parent + 1,000 tenants = ~2.5MB/tenant

Approach	Storage	Cost
Full Copy	512GB	\$51,200/mo
RVCOW	~2.5GB	\$250/mo
Savings	200x	99.5%

COW branch creation: **2.6ms**

CowMap cluster lookup: **28ns**

Row-level security per tenant

61μs

HNSW Search

8,000x

vs Pinecone

290+

SQL Functions

200x

COW Savings

RuVector isn't just faster search. It's a fundamentally new data architecture that unlocks capabilities impossible with traditional vector databases — in-process execution, WASM portability, and native PostgreSQL integration that eliminates the 'AI sidecar' pattern entirely.

RVF: The Executable Knowledge Unit

24-segment binary format. Not a database file — an executable that contains intelligence.

Core Data (3)

Vectors, Metadata, Indexes

AI / Model (3)

LoRA deltas, GGUF weights, Inference config

COW Branching (4)

Git-like versioning (200x savings)

Security (2)

Witness chains (ML-DSA-65), TEE attestations

Compute (3+)

WASM (5.5KB) | eBPF | Unikernel (125ms)

AGI Container (9)

Identity + tools + authority + evaluation

Docker image: 50-500 MB

Minimal container: 5-50 MB

RVF WASM runtime: 5.5 KB

Docker images are megabytes. An RVF WASM runtime is 5.5 kilobytes.

AIMDS: Security via Chaos Theory

3-layer pipeline. Lyapunov chaos detection. Mathematical certainty, not heuristics.

STAGE 1: DETECTION

0.06ms

- 50+ detection patterns
- Aho-Corasick matching
- Real-time stream processing

STAGE 2: ANALYSIS

80ms

- Lyapunov chaos detection
- LTL policy verification
- Behavioral divergence

STAGE 3: RESPONSE

<50ms

- 7 adaptive strategies
- 25-level strange-loop
- Real-time mitigation

Lyapunov Chaos Detection

Positive Lyapunov exponent = system diverging from expected behavior = adversarial. Same mathematics that predicts weather chaos. Detects novel attacks without pattern libraries.

782KB

Binary Size

>12K/s

Requests/Sec

<5%

False Positive

0.06ms

Detection Time

SONA + ReasoningBank: Real-Time Learning

Every decision makes the system smarter. Automatically. <1ms. \$0.

SONA: Two-Tier LoRA

- MicroLoRA (rank-2, ~45 μ s): instant reactions
- BaseLoRA (rank-16, ~500 μ s): deep learning
- \$0 cost, <1ms latency, zero downtime
- EWC++ prevents catastrophic forgetting (45% less)

Traditional Fine-Tuning

- Days to weeks of training time
- \$1K-\$100K per fine-tune cycle
- Requires scheduled downtime
- Catastrophic forgetting on new domains

ReasoningBank 4-Step Pipeline

1. RETRIEVE

HNSW search for similar past decisions

2. JUDGE

Assign verdicts based on outcomes

3. DISTILL

Micro-LoRA pattern extraction

4. CONSOLIDATE

EWC++ prevents forgetting

300x

Faster Retrieval

2,211/s

SIMD Ops/Sec

+55%

Quality Improvement

\$0

Learning Cost

Ship It In 5 Lines

Production multi-agent swarms. Self-learning. Self-healing. One require statement.

```
const ruflo = require('ruflo');

// Initialize a self-healing agent swarm
const swarm = await ruflo.swarm.init({
  topology: 'hierarchical', agents: 8, strategy: 'specialized'
});

// Deploy 51 autonomous agents in one command
const fleet = await swarm.deploy({
  domains: ['billing', 'compliance', 'support'],
  memory: 'persistent',           // AgentDB - agents remember everything
  security: 'aimds',              // Chaos theory threat detection
  runtime: 'rvf-wasm'             // 5.5KB - runs anywhere
});

fleet.on('resolution', (data) => {
  // SONA generates LoRA adapters automatically
  // Your AI gets smarter. Zero cost. Zero effort.
});
```

Claude Code shipped sub-agents in March 2026. Ruflo shipped this in August 2025. 8 months earlier.

Bio-Inspired Nervous System

Not a metaphor — implemented computational neuroscience. 22.9K lines, 359 tests.

Layer 1: SENSING <100ns

Lock-free ring buffers, 10K+ events/ms

Layer 2: REFLEX <1μs

K-Winner-Take-All, dendritic coincidence detection

Layer 3: MEMORY 10^{30} cap

Hyperdimensional Computing, XOR binding <50ns

Layer 4: LEARNING $O(1)$ /synapse

BTSP one-shot, E-prop online, EWC++ anti-forgetting


Layer 5: COHERENCE 40Hz gamma

Kuramoto oscillators, global workspace, predictive coding

Circadian Controller: 5-50x compute savings during quiet periods — the system sleeps when you do

We Built the Future 8 Months Early

RuvNet doesn't follow the industry roadmap. The industry follows ours.



Two standard deviations beyond state of the art. Open source. MIT license. Free.

Graph Analytics: MinCut + GNN + HNSW

Stoer-Wagner algorithm. 448+ tests. Production graph intelligence.

MinCut Applications

- Finding architectural bottlenecks
- Identifying weak points in distributed systems
- Detecting natural module boundaries
- Measuring attack surface
- Swarm topology optimization
- Anti-pattern detection via absence analysis

GNN on HNSW Topology

- GCN, GAT, GraphSAGE implementations
- GraphRoPE: hop-distance-aware embeddings
- Louvain community detection
- Spectral clustering (K-way)
- Bridge/articulation point detection
- Dynamic MinCut for real-time health

Industry Proof Points

Pinterest

150% hit-rate improvement
Graph-based recommendation

Google Maps


50% better ETA
Graph routing optimization

Uber Eats

20%+ engagement
Graph-ranked discovery

Performance Benchmarks

Measured results against best-in-class competitors.



Metric	RuVector	Best Competitor	Advantage
HNSW Search	61μs	500ms+ (Pinecone)	8,000x
Fox Flow QPS	12.8M	3K (Redis)	4,000x
Detection Time	0.06ms	~10ms (rule-based)	167x
WASM Runtime	5.5KB	100MB+ (containers)	18,000x
Learning Latency	<1ms	Days-weeks	>86Mx
Reflex Response	<1μs	N/A	Unique
Agent Scale	100K	~10-50	2,000x+
SWE-Bench	84.8%	33.2% (GPT-4o)	2.5x
Attention Mechanisms	39 types	1-3 types	13-39x

8,000x

Vector Search

4,000x

Stream Processing

167x

Threat Detection

18,000x

Smaller Runtime

Integration Architecture

MCP protocol. PostgreSQL-native. Drop-in pgvector replacement.

MCP Protocol: 96 Tools

- Deferred loading for performance
- Agent orchestration + memory + search
- Task management + workflow engine
- Security scanning + graph analytics

PostgreSQL-Native

- No separate vector database deployment
- 290+ SQL functions in-process
- ACID transactions on vector ops
- Drop-in pgvector migration

APIs, SDKs & Integration Points

REST API

Standard HTTP

gRPC

High-perf binary

TypeScript SDK

Full type safety

Python SDK

ML pipelines

WASM SDK

Client-side search

Webhooks

Event-driven

Implementation Roadmap

From zero to production in 3 months. Incremental value at every phase.

1: Foundation

Week 1-2

- Core RuVector PostgreSQL extension
- AIMDS security middleware
- Basic agent deployment
- HNSW index configuration

2: Intelligence

Week 3-4

- AgentDB memory (4 types)
- ReasoningBank + SONA
- Ruflo orchestration
- MCP tool integration

3: Customize

Month 2

- Custom agent development
- Domain-specific LoRA training
- Graph analytics deployment
- Performance tuning

4: Scale

Month 3+

- Production scaling + monitoring
- WASM edge deployment
- Multi-tenant RVCOW
- Advanced swarm topologies

Each phase delivers standalone value. No big-bang migration required.

R U V N E T

Schedule a 30-Minute Architecture Review

148+

Capabilities

14

Modules

80+

Rust Crates

290+

SQL Functions

39

Attention Mechanisms

9

RL Algorithms

5

Bio-Inspired Layers

24

RVF Segment Types

Two standard deviations beyond state of the art. Open source. Free. The architecture everyone will use.

ruvnet.com | hello@ruvnet.com